ACHELION CAPITAL MANAGEMENT, LLC

# ARMS Intelligence Architecture

## Phase 2 and Phase 3 — Anticipatory Signal Layer + Full Execution Autonomy

For: Josh Paul — Lead Developer

Version: 1.0 — March 26, 2026 · Addendum 3 to FSD v1.1

**Classification: CONFIDENTIAL — GP and Development Use Only**

**The PM staring at a monitor is the failure mode. This document specifies the architecture that puts the PM in the room where Category A intelligence is created — not watching a dashboard waiting to approve what the system already knows.**

**"See it before the internet does."**

# 0. The Design Principle That Governs Everything

> *Every minute the PM spends confirming what the system already knows is a minute not spent in the room where Category A intelligence is created. The architecture exists to free the PM entirely from the execution and decision layer — so that the only thing the PM does is see what the data cannot yet see.*

This document specifies the architecture that operationalizes that principle completely. It is built on a single governing question that Josh must ask about every human touchpoint in the system:

**Is this decision based on math, process, and data inflow — or does it require pre-consensus human intelligence that no data feed can provide?**

If the answer is math, process, or data — it is automated. No exceptions. No confirmation windows for deterministic outcomes. No decision queue items that repeat daily without change. The system executes. The PM is not notified unless the system fails.

If the answer is pre-consensus human intelligence — it goes to the PM immediately, with full context, structured as a specific question that requires a specific answer. Not a general alert. Not a price notification. A structured question: does your Category A knowledge change this recommendation?

The architecture that follows is the complete implementation of this principle across the anticipatory intelligence layer and the full execution autonomy layer. Phase 2 builds the intelligence. Phase 3 closes the last human touchpoints in execution. By the end of Phase 3, the PM's interaction with ARMS is three things: seeding new thesis ideas through SENTINEL, answering specific Category A questions when the system surfaces them, and reviewing the monthly session log analytics. Everything else runs without them.

# 1. The Three Layers of Anticipatory Intelligence

Anticipatory intelligence operates at three depths of forward visibility. Each layer sees further ahead than the one before it. Each requires more sophisticated signal processing. All three are buildable with the architecture already specified. Josh builds them in sequence.

| Layer | Name | Signal type | Forward visibility | Phase |
|---|---|---|---|---|
| **Layer 1** | **Reactive speed** | Legal rulings, regulatory filings, earnings warnings, insider sales — events already on the internet | 30 minutes ahead of human reading speed. Market may already be moving. | **Phase 1 — CDM + TDC already specified** |
| **Layer 2** | **Pattern detection before consensus** | Earnings language velocity, job posting trends, patent filing velocity — public data not yet connected by consensus | 1 to 3 quarters ahead of market consensus framing | **Phase 2 — this document** |
| **Layer 3** | **Upstream condition detection** | Supply chain manifests, satellite data, cross-industry capital flow patterns, geopolitical preconditions | 2 to 4 quarters ahead — the conditions that precede the conditions | **Phase 3 — scales with AUM and data budget** |

# 2. Layer 2 — Anticipatory Signal Specifications

**These four signals are the pre-consensus intelligence layer. All are free. All are buildable in Phase 2. All connect directly to the CDM thesis dependency structure already built.**

## 2.1 Earnings Language Velocity Tracker (ELVT)

| | |
|---|---|
| **File** | src/intelligence/elvt.py — new module |
| **Data source** | SEC EDGAR full-text search — earnings call transcripts via 8-K filings. Free API. |
| **Runs** | Quarterly after each earnings season + on-demand when CDM fires for a named entity |
| **Build time** | 1 week — Claude API already specified. Transcript parsing is incremental. |
| **Cost** | ~$0.30 per transcript analysis via Claude API. ~$30-50/month at full Architecture AB coverage. |

ELVT tracks how language changes across consecutive earnings calls for every named entity in the CDM map. Not what they say — how they say it, and whether the direction is changing. Three language categories are scored on every transcript.

| Language category | What it measures | Bullish signal | Bearish signal |
|---|---|---|---|
| Capex commitment | Are they saying 'we will invest' or 'we are evaluating'? Commitment language vs optionality language. | Shift from conditional to definitive capex language across 2+ consecutive calls | Shift from definitive to conditional or absent capex language |
| Demand visibility | Are they describing backlog growth, customer urgency, lead time extensions? Or cautious ordering, inventory normalization? | Lead time extension language + customer urgency + backlog expansion across 2+ calls | Inventory normalization + order push-out language + softening urgency |
| Competitive posture | Are they describing pricing power, design wins, sole-source positions? Or pricing pressure, competitive alternatives, share loss? | Pricing power + new design wins + expanding TAM language | Pricing pressure + competitive alternatives mentioned + TAM language softening |

The ELVT score is not the content of any single transcript. It is the velocity of change across consecutive transcripts. A company that shifts from conditional to definitive capex language across three consecutive quarters is describing an investment decision that has been made internally but has not yet appeared in a financial filing. That is a 2-3 quarter leading indicator.

### Current book application — the signals that matter most

| Position | Named entities to track | Signal that matters | Forward indicator of |
|---|---|---|---|
| MU | Microsoft, Google, Amazon, Meta earnings calls | Capex commitment language shifting to definitive on AI infrastructure spend | HBM demand acceleration 2-3 quarters out |
| NVDA | Microsoft, Google, Amazon, Meta, Oracle | Demand visibility: lead time extension language + urgency language | GPU allocation tightening — revenue acceleration |
| ALAB | Microsoft, Google, Amazon, Meta | PCIe/CXL connectivity urgency language in data center discussions | ALAB design win acceleration |
| AVGO | Google, Apple, Meta | Custom ASIC design language — partner references without naming | New ASIC customer pipeline |
| PLTR | DOD press releases, Congressional budget hearings | AI program urgency language + sole-source contract language | Contract expansion cycle |
| ANET | Microsoft, Meta, Google | Network upgrade urgency language + capacity constraint language | AI networking demand cycle |

## 2.2 Job Posting Velocity Intelligence (JPVI)

| | |
|---|---|
| **File** | src/intelligence/jpvi.py — new module |
| **Data source** | Adzuna API (already in FSD v1.1 free signal feeds) + LinkedIn public job data |
| **Runs** | Weekly — same Monday cadence as Systematic Scan Engine |
| **Build time** | 3 days — Adzuna already connected. Adds structured query framework against CDM map. |
| **Cost** | $0 — Adzuna free tier. LinkedIn scraping via public postings. |

When a hyperscaler is about to accelerate data center investment, its job postings in specific categories increase 6 to 9 months before the capital spend appears in any financial filing. This is because hiring decisions precede procurement decisions — you hire the team before you buy the equipment. JPVI tracks posting velocity for named CDM entities across four job categories that are leading indicators of AI infrastructure spend.

| Job category | Why it matters | Lead time | Positions it affects |
|---|---|---|---|
| Data center infrastructure engineering | These roles build the physical layer that buys AI chips and memory | 6-9 months before capex | MU, NVDA, ALAB, ANET |
| AI/ML infrastructure and platform | These roles deploy the training and inference systems that consume GPU and HBM | 4-6 months before procurement | NVDA, MU, AVGO, MRVL |
| Network engineering — AI cluster connectivity | These roles specify the networking layer for AI clusters | 4-6 months before networking procurement | ANET, MRVL, AVGO |
| Power and cooling engineering | These roles build the infrastructure that enables AI data center density expansion | 6-12 months before power procurement | ETN, VRT |

A 30% month-over-month increase in data center infrastructure engineering postings at Microsoft or Google is not a noise signal. It is a company that has decided to build and is hiring to execute. The equipment orders follow 6 to 9 months later. ARMS sees the hiring before the market sees the orders.

### JPVI velocity scoring

```
def score_jpvi_velocity(entity: str, category: str,
                        current_30d: int,
                        prior_90d_avg: float) -> JPVISignal:

    velocity = (current_30d - prior_90d_avg) / prior_90d_avg

    if velocity > 0.40:   return JPVISignal('ACCELERATING', velocity, 'CRITICAL')
    if velocity > 0.20:   return JPVISignal('INCREASING',   velocity, 'HIGH')
    if velocity > 0.05:   return JPVISignal('ELEVATED',     velocity, 'MEDIUM')
    if velocity < -0.20:  return JPVISignal('DECLINING',    velocity, 'HIGH')
    return JPVISignal('STABLE', velocity, 'LOW')
```

## 2.3 Patent Filing Velocity Tracker (PFVT)

| | |
|---|---|
| **File** | src/intelligence/pfvt.py — new module |
| **Data source** | USPTO PatentsView API — free bulk data. Updated weekly. |
| **Runs** | Monthly — lower frequency signal. Patent velocity changes slowly but leads by 12-36 months. |
| **Build time** | 1 week — USPTO API is documented and free. Mapping categories to thesis structures is the primary design work. |

| | |
|---|---|
| **Cost** | $0 — USPTO bulk data is a public good. |

Patent applications are filed 12 to 18 months before products ship and 24 to 36 months before they reach material revenue. When a company begins filing patents in a specific technology category, it is describing its competitive position 2 to 3 years from now. When a hyperscaler begins filing patents in a new memory architecture or networking topology, it is describing its infrastructure requirements 2 years hence — requirements that named ARMS positions will be asked to fulfill.

| Patent category | Monitored entities | Lead time | Thesis implication |
|---|---|---|---|
| HBM and advanced memory packaging | Google, Microsoft, Amazon, Samsung, SK Hynix, Micron | 24-36 months | HBM demand cycle shape and next-generation spec requirements for MU |
| AI inference chip architecture | Google (TPU), Microsoft, Amazon (Trainium/Inferentia), Meta | 18-30 months | Custom ASIC vs NVDA GPU substitution risk. Affects NVDA conviction score. |
| PCIe Gen 6 and CXL interconnect | Intel, AMD, ALAB, hyperscalers | 18-24 months | ALAB design win pipeline — PCIe Gen 6 adoption velocity |
| Data center power density and cooling | Google, Microsoft, Amazon, ETN, VRT | 24-36 months | Next-generation data center power requirements. ETN and VRT total addressable market. |
| Custom networking silicon | Microsoft, Meta, Google, MRVL, AVGO | 18-30 months | Custom networking adoption vs merchant silicon — affects ANET, MRVL, AVGO |

## 2.4 Supply Chain Cross-Reference (SCCR)

| | |
|---|---|
| **File** | src/intelligence/sccr.py — new module |
| **Data source** | Phase 2: MarineTraffic free tier + US import/export manifest data (free via ImportYeti/Panjiva open data). Phase 3: Orbital Insight satellite data (institutional cost). |
| **Runs** | Weekly — shipping manifest data updates with 2-4 week lag but patterns emerge over monthly windows. |
| **Build time** | 2 weeks — data parsing is more complex than other feeds. Phase 2 free tier is sufficient for pattern detection. |
| **Cost** | Phase 2: $0. Phase 3: institutional satellite data pricing. |

Physical capital movement precedes financial reporting by 2 to 4 quarters. When TSMC ships an unusually large volume of advanced packaging materials, something significant is being built. When hyperscaler equipment orders arrive at specific ports in unusual volume, a data center expansion is

underway that will not appear in any earnings call for 6 months. The supply chain does not lie — it is the physical reality of investment decisions already made.

Phase 2 SCCR monitors two specific patterns against the CDM dependency map. The first is TSMC-related shipping activity from Taiwan to major US and Asian data center hubs. The second is hyperscaler server equipment import volumes at key US ports. Both are available in public manifest data. Both provide 2 to 4 quarter leading indicators of the AI infrastructure demand that flows to NVDA, MU, ALAB, ANET, AVGO, and MRVL.

# 3. The Closed Learning Loop — How ARMS Gets Smarter

**The signals are not the edge. The edge is having a thesis framework that tells the system which signals matter — and a learning loop that continuously improves the signal weights based on which ones were actually predictive.**

The four Layer 2 signals produce data. The learning loop converts that data into improving signal weights over time. This is the machine learning layer that Renaissance built over 30 years. ARMS builds it from day one — because the session log, the MICS formula, and the CCM are the architecture of the learning loop.

The loop runs quarterly. It asks three questions about each signal across each thesis.

| Question | Method | Output |
|---|---|---|
| Was this signal predictive? | Correlate signal readings from 3-6 months ago against subsequent thesis performance (TIS score movement, price performance vs QQQ, MICS score accuracy). Spearman rank correlation. | Predictive accuracy score per signal per thesis type. Stored in session log. |
| Was the signal lead time correct? | Measure the actual lag between signal elevation and thesis outcome. Compare to assumed lead time. Adjust lead time assumption. | Updated lead time estimates per signal per entity. Feeds JPVI and ELVT timing. |
| Did any signal combination outperform any single signal? | Test pairwise and three-way signal combinations against outcomes. Identify whether ELVT + JPVI together is more predictive than either alone. | Signal combination weights for multi-signal scoring. Feeds forward into MICS Gate 3 supplementary scoring. |

The output of the learning loop is not a trading signal. It is a weight update — which signals get more influence in the Gate 3 mispricing score and which get less, based on their actual track record of predicting thesis outcomes for each specific position type. Over 8 to 12 quarters of data, the system builds a position-type-specific signal model that is continuously validated against its own history.

This is the exact architecture Renaissance used. They did not know in advance which signals would be predictive. They built the infrastructure to measure it and let the data decide. ARMS does the same — with the additional constraint that every signal must have a causal explanation rooted in the SENTINEL thesis framework. No correlation without causation. That constraint is what prevents overfitting.

# 4. Full Execution Autonomy — Closing Every Human Touchpoint

**This section is the most important in the document. It specifies, exhaustively, every remaining human touchpoint in ARMS and the specific architecture that eliminates each one.**

The goal is a system where the PM's daily interaction with ARMS is zero unless one of three conditions is true: (1) a Category A intelligence question has been surfaced that requires pre-consensus human judgment, (2) a SENTINEL thesis seed requires the PM's Gate 6 source declaration, or (3) a GP-level structural decision has been triggered. Everything else executes without PM involvement.

The following table maps every current human touchpoint to its elimination specification. Josh works through this list systematically. Each row is a build task.

| Current human touchpoint | Current tier | Target tier | Elimination specification |
|---|---|---|---|
| PTRH sizing confirmation | Tier 1 | Tier 0 | Module A (PTRH Automation) — already specified. Build first. |
| CDF weight decay application | Tier 1 | Tier 0 | CDF fires automatically when day count and performance threshold are met. No veto window. No confirmation. System reduces weight and logs. PM sees result in module status, not decision queue. |
| DSHP defensive sleeve harvest | Tier 1 | Tier 0 | Module B (DSHP) currently specified as Tier 1. After 90 days of validated operation with zero PM vetoes, promote to Tier 0. PM sets the thresholds once. System harvests automatically. |
| ARES re-entry tranche execution | Tier 1 | Tier 0 | ARES gates are deterministic. When all 4 gates clear and VARES sizes the tranche, the order executes. No veto window. The PM has already declared the thesis. Regime clearing is the execution trigger. |
| PERM covered call evaluation | Tier 1 | Tier 0 | When unrealized gain exceeds 30% AND the options market is open AND IV is above a defined threshold (VIX-calibrated), the covered call is executed automatically. Strike selection: ATM. No PM confirmation. |
| TDC thesis review PM response | Tier 1 | Tier 1 — preserve | This is the correct human touchpoint. When TIS falls to IMPAIRED or BROKEN, the PM answers one question: do you have |

| | | | |
|---|---|---|---|
| | | | Category A knowledge that explains why the thesis remains intact despite the evidence? If yes, they provide it. If silence, system proceeds with CDF acceleration. |
| SENTINEL new position initiation | Tier 1 | Tier 1 — preserve | This is the correct human touchpoint. The PM seeds the thesis. The system runs all 6 gates. If all pass, the position initiates. PM can provide Category A context via Gate 6. Silence = system executes at MICS-implied sizing. |
| FEM ALERT trim recommendation | Tier 1 | Tier 0 | When FEM reaches ALERT (>80%) for more than 24 hours AND a specific trim recommendation is identified by the Master Engine AND the trim does not require a full SENTINEL re-run, it executes automatically. Paired trim is deterministic. |
| Regime ceiling enforcement | Tier 0 | Tier 0 — already correct | No change. Ceiling is enforced immediately on regime transition. Always has been. |
| PDS DELEVERAGE execution | Tier 0 | Tier 0 — already correct | No change. Drawdown triggers fire immediately. Always have. |
| MICS score assignment | Currently PM-assigned | Tier 0 — MICS formula | When MICS module is live, conviction scores are formula-derived. PM override is ±1 level with written reason. This is already specified in FSD v1.1. Build priority #1. |
| Thesis Retirement Protocol | Tier 1 | Tier 0 after CCM live | Once CCM has 90 days of data, TRP can execute automatically when TIS score is BROKEN for 30+ consecutive days AND CDF is at 0.60× AND no PM conviction reaffirmation has been logged in 45 days. Orderly exit executes. |
| VRT PERM check | Recurring manual | Tier 0 | Captured by PERM Tier 0 promotion above. Covered call evaluation triggers when unrealized gain crosses 30%. No manual check required. |
| GOOG regime queue monitoring | PM monitoring | Tier 0 | ARES and regime module monitor this automatically. When regime clears to NEUTRAL, GOOG initiation fires without PM involvement. PM declared the thesis. System executes the entry. |

## 4.1 The Two Legitimate Human Touchpoints — Defined and Bounded

After full execution autonomy is achieved, exactly two human touchpoints remain in the ARMS architecture. Both are legitimate. Neither is eliminatable. Both are the PM doing what only the PM can do.

| Touchpoint | What it requires | What makes it irreplaceable |
|---|---|---|
| **SENTINEL thesis seeding (Gate 6 source declaration)** | PM identifies a company representing a civilizational shift before it is in any dataset. PM declares the source of insight: direct conversation (Cat A), pattern recognition (Cat B), or public synthesis (Cat C). System runs gates 1-5 automatically. Gate 6 is the only human gate. | No algorithm can replicate 28 years of market pattern recognition applied to identifying what the world will look like before the data confirms it. This is the PM's irreducible contribution. Everything else is the system's job. |
| **Category A intelligence response (TDC IMPAIRED/BROKEN flag)** | When TDC scores a position as IMPAIRED or BROKEN, it surfaces one specific question: do you have Category A knowledge — a direct conversation, a pattern recognition from prior cycles — that explains why the thesis remains intact despite the evidence? PM responds or silently confirms. | The system knows everything that is knowable from public data. It cannot know what the PM learned at dinner last week. The Category A response gate is the mechanism by which private pre-consensus knowledge enters the system in a documented, auditable way. |

> *Two touchpoints. Both irreplaceable. Everything else automated. The PM is in Dubai — not staring at a monitor.*

# 5. Complete Phase 2 and Phase 3 Build Plan for Josh

## 5.1 Phase 2 Build Sequence (Months 3-4)

**Phase 2 goal: anticipatory signal layer live. Every execution touchpoint promoted to Tier 0 except the two legitimate human gates.**

| # | Build item | Time | Cost | Milestone |
|---|---|---|---|---|
| 1 | CDF Tier 0 promotion — remove veto window, execute automatically | 1 day | $0 | CDF never appears in decision queue. Weight reduction is instantaneous and logged. |
| 2 | PERM Tier 0 promotion — auto-execute covered call when conditions met | 2 days | $0 | Covered call executes automatically on >30% gain + VIX threshold. VRT case resolves permanently. |
| 3 | ARES Tier 0 promotion — auto-execute re-entry tranches | 2 days | $0 | Re-entry executes when all 4 gates clear. PM not notified unless gate failure. |
| 4 | FEM paired trim Tier 0 — auto-execute when ALERT + trim identified | 2 days | $0 | FEM ALERT no longer a recurring decision queue item. |
| 5 | ELVT — earnings language velocity tracker | 1 week | ~$30-50/month | Quarterly earnings calls for all CDM entities scored automatically. Language velocity signals feeding Gate 3 supplementary score. |
| 6 | JPVI — job posting velocity intelligence | 3 days | $0 | Weekly hyperscaler hiring trend signals feeding anticipatory layer. 6-9 month forward visibility on AI capex. |
| 7 | PFVT — patent filing velocity tracker | 1 week | $0 | Monthly patent filing trends for named CDM entities. 12-36 month thesis visibility. |
| 8 | SCCR — supply chain cross-reference (free tier) | 2 weeks | $0 | Weekly shipping manifest patterns for TSMC and hyperscaler equipment. 2-4 quarter leading indicators. |
| 9 | Learning loop v1 — quarterly signal calibration | 1 week | $0 | First quarterly CCM + signal weight recalibration run. System begins self-improving. |

| | | | | |
|---|---|---|---|---|
| 10 | DSHP Tier 0 promotion (after 90-day validation) | 1 day | $0 | Defensive sleeve harvest executes automatically. No veto window. |

## 5.2 Phase 3 Build Sequence (Months 5-6)

**Phase 3 goal: full FSD achieved. TRP autonomous. LP digest automated. ARMS is a licensable platform.**

| # | Build item | Time | Cost | Milestone |
|---|---|---|---|---|
| 1 | TRP Tier 0 promotion (after CCM live + 90 days data) | 2 days | $0 | Thesis Retirement Protocol executes automatically when 3 conditions met. Orderly exit without PM involvement. |
| 2 | SCCR upgrade — satellite data integration | 1 week | AUM-dependent | Phase 3 upstream condition detection. 2-4 quarter forward visibility on physical infrastructure build. |
| 3 | Proactive Intelligence Digest (PID) — monthly LP report | 1 week | $0 (Claude API already budgeted) | Automated monthly LP report delivered without PM involvement. Transparent audit of every regime decision. |
| 4 | ARMS licensing config layer | 2 weeks | $0 | Fund-agnostic parameter set. ARMS becomes a platform that can be licensed to other funds. |
| 5 | Performance attribution by module | 1 week | $0 | Every basis point of alpha attributed to the specific ARMS module that generated it. LP-defensible track record documentation. |
| 6 | Learning loop v2 — signal combination optimization | 2 weeks | $0 | Multi-signal combination testing. System identifies which signal pairs are more predictive than individual signals. |

# 6. What the PM's Day Looks Like After Full Build

**This is the design target. Every build decision Josh makes should be measured against whether it moves closer to this picture.**

| Scenario | What happens | PM involvement |
|---|---|---|
| Normal market day — no signals | ARMS runs all modules. Monitor generated at 6AM CT. Portfolio operating within regime ceiling. No TDC flags. No CDM alerts. | Zero. PM wakes up to a monitor that confirms everything is running correctly. No action required. |
| Regime transition — DEFENSIVE to NEUTRAL | Macro Compass score crosses 0.65. ARAS immediately enforces new ceiling. ARES begins evaluating re-entry gates. PTRH resizes to 1.25×. SLOF eligibility checked. GOOG initiation queued and executed if all gates clear. | Zero. PM reads the monitor the next morning and sees the transition executed correctly. No confirmation required. |
| Named entity legal event — Google antitrust ruling | CDM fires within 30 minutes. 7 positions receive cross-reference alerts. TDC runs thesis reviews for all 7. Results logged. WATCH flags updated on affected positions. Monitor reflects updated status. | Zero for reactive speed. PM receives a summary in the monitor. If TDC scores any position IMPAIRED, PM receives one specific question: do you have Category A knowledge that changes this assessment? |
| JPVI acceleration signal — Microsoft hiring surge | JPVI detects 45% month-over-month increase in Azure data center engineering postings. Signal scored ACCELERATING. Feeds Gate 3 supplementary score for MU, NVDA, ALAB, ANET. MICS scores recalculated. | Zero. If MICS scores cross a threshold that changes conviction weights materially, PM sees the summary in the monitor. No confirmation required unless change triggers FEM ALERT. |
| PERM triggered — VRT crosses 30% gain | Covered call evaluation executes automatically. IV check passes (VIX above threshold). ATM covered call sold. Premium logged. Session log updated. | Zero. PM sees confirmation in module status: 'PERM executed: VRT covered call sold [date] [premium]. Next evaluation in 30 days.' |
| TDC IMPAIRED — position thesis weakening | TDC scores a position IMPAIRED. System surfaces one question to PM via mobile push: 'TDC has scored [POSITION] thesis as IMPAIRED. Bear case: [specific evidence]. Do you have | One response. PM answers yes (and provides the knowledge) or no (and system proceeds with CDF acceleration). Total time: 5 minutes. |

| | | |
|---|---|---|
| | Category A knowledge that explains why the thesis remains intact?' | |
| SENTINEL seed — PM identifies new thesis | PM is at dinner in Dubai. Hears something from a TSMC supply chain contact that suggests a specific company is about to become non-optional AI infrastructure. PM opens the SENTINEL interface on their phone and seeds the thesis with a Gate 6 Cat A declaration. | One action — thesis seeding. System runs all gates automatically overnight. PM wakes up to a complete SENTINEL analysis. If all gates pass, PM has a 4-hour veto window (default execute). Silence = system initiates. |
| Monthly learning loop | CCM runs quarterly signal calibration. MICS formula weights updated. Signal lead time estimates refined. Session log analytics delivered. | Zero. PM reviews the monthly analytics digest as strategic context. No operational decisions required. |

> *The PM is in Dubai. The system is running. The portfolio is being managed. The only time the phone buzzes is when the system has a question only the PM can answer.*

# 7. The Architecture in One Statement

The goal of ARMS is not to build a better dashboard. It is not to give the PM better information faster. It is to remove the PM from every decision that can be made by math, process, and data — so that the PM's entire cognitive capacity is reserved for the one thing that produces irreplaceable alpha: seeing what the data cannot yet see.

The anticipatory signal layer described in this document is the system seeing as far ahead as public data allows. The two legitimate human touchpoints are the system asking the PM to see further still — through the conversations, the pattern recognition, and the pre-consensus intelligence that 28 years of market experience produces and no algorithm can replicate.

When this architecture is fully built, Achelion will operate with a level of systematic discipline, anticipatory intelligence, and execution autonomy that no emerging fund — and very few institutional funds — can match. The edge is real. The architecture is right. The build sequence is clear.

**See it before the internet does. Act before the market does. Let the system do everything the system can do. Let the PM do everything only the PM can do.**

*"The PM is in Dubai. The system is running."*

Achelion Capital Management, LLC · Flow. Illumination. Discipline. Conviction.